

# Using AmgX to Accelerate PETSc-Based CFD Codes

Pi-Yueh Chuang

[pychuang@gwu.edu](mailto:pychuang@gwu.edu)

George Washington University

# Our Group

- Professor Lorena A. Barba  
<http://lorenabarba.com/>
- Projects:
  - **PyGBe** - *Python GPU code for Boundary elements*  
<https://github.com/barbagroup/pygbe>
  - **PetIBM** - *A PETSc-based Immersed Boundary Method code*  
<https://github.com/barbagroup/PetIBM>
  - **cuIBM** - *A GPU-based Immersed Boundary Method code*  
<https://github.com/barbagroup/cuIBM>
  - ... *and so on*  
<https://github.com/barbagroup>

# **Our story**

How we painlessly enable multi-GPU computing in PetIBM

# PETSc

- Portable, Extensible Toolkit for Scientific Computation  
<https://www.mcs.anl.gov/petsc/index.html>
- Argonne National Laboratory, since 1991
- Intended for large-scale parallel applications
- Parallel vectors, matrices, preconditioners, linear & nonlinear solvers, grid and mesh data structure ... etc
- Hides MPI from application programmers
- C/C++, Fortran, Python

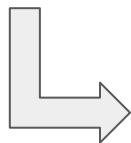
# PetIBM

Taira & Colonius' method (2007):

$$\frac{\partial \mathbf{u}}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{u} = -\nabla p + \frac{1}{Re} \nabla^2 \mathbf{u} + \int_s \mathbf{f}(\boldsymbol{\xi}(s, t)) \delta(\boldsymbol{\xi} - \mathbf{x}) ds,$$

$$\nabla \cdot \mathbf{u} = 0,$$

$$\mathbf{u}(\boldsymbol{\xi}(s, t)) = \int_{\mathbf{x}} \mathbf{u}(\mathbf{x}) \delta(\mathbf{x} - \boldsymbol{\xi}) d\mathbf{x} = \mathbf{u}_B(\boldsymbol{\xi}(s, t)),$$



$$\begin{bmatrix} A & G & E^T \\ G^T & 0 & 0 \\ E & 0 & 0 \end{bmatrix} \begin{pmatrix} q^{n+1} \\ \phi \\ \tilde{f} \end{pmatrix} = \begin{pmatrix} r^n \\ 0 \\ u_B^{n+1} \end{pmatrix} + \begin{pmatrix} bc_1 \\ -bc_2 \\ 0 \end{pmatrix}$$

<sup>†</sup>K. Taira and T. Colonius, "The immersed boundary method: A projection approach", Journal of Computational Physics, vol. 225, no. 2, pp. 2118-2137, 2007.

# PetIBM

$$Q \equiv [G, \ E^T], \quad \lambda \equiv \begin{pmatrix} \phi \\ \tilde{f} \end{pmatrix}, \quad r_1 \equiv r^n + bc_1, \quad r_2 \equiv \begin{pmatrix} -bc_2 \\ u_B^{n+1} \end{pmatrix}$$



$$\begin{bmatrix} A & 0 \\ Q^T & -Q^T B^N Q \end{bmatrix} \begin{bmatrix} I & B^N Q \\ 0 & I \end{bmatrix} \begin{pmatrix} q^{n+1} \\ \lambda \end{pmatrix} = \begin{pmatrix} r_1 \\ r_2 \end{pmatrix} + \begin{pmatrix} -\frac{\Delta t^N}{2^N} (LM^{-1})^N Q \lambda \\ 0 \end{pmatrix}$$



$$Aq^* = r_1 \quad (\text{Solve for intermediate velocity}),$$

$$Q^T B^N Q \lambda = Q^T q^* - r_2 \quad (\text{Solve the modified Poisson equation}),$$

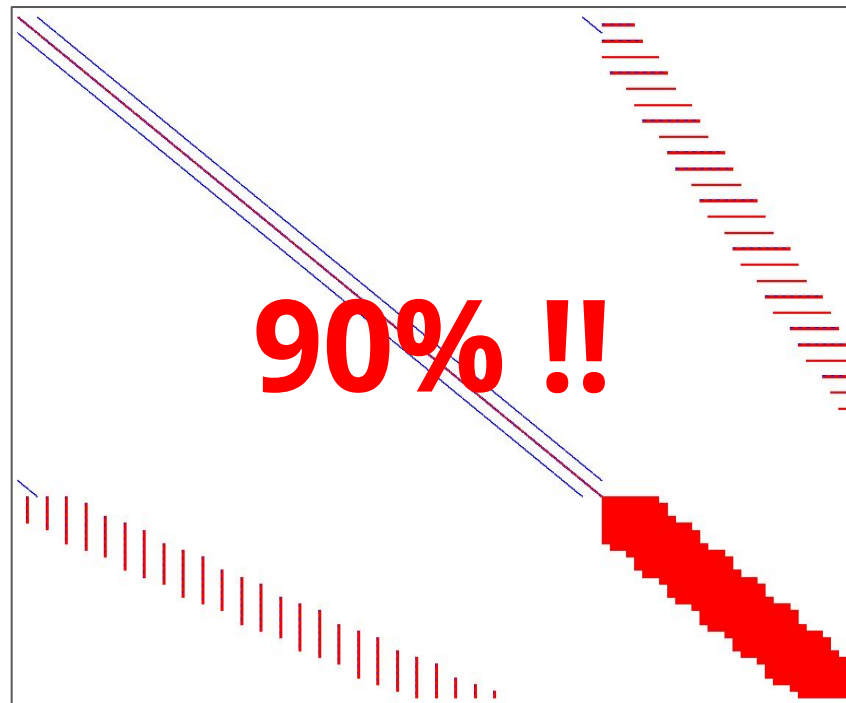
$$q^{n+1} = q^* - B^N Q \lambda \quad (\text{Projection step}).$$

# Solving modified Poisson systems is tough

## Possible solutions:

Rewrite the whole program for  
multi-GPU capability, or

Tackle the expensive part !



# AmgX

- **Developed and supported by NVIDIA**  
<https://developer.nvidia.com/amgx>
- **Krylov methods:**
  - CG, GMRES, BiCGStab, ... etc
- **Multigrid preconditioners:**
  - Classical AMG (largely based on Hypra BoomerAMG)
  - Unsmoothed aggregation AMG
- **Multiple GPUs on single node / multiple nodes:**
  - MPI (OpenMPI) / MPI Direct
  - Single MPI rank  $\Leftrightarrow$  single GPU
  - Multiple MPI ranks  $\Leftrightarrow$  single GPU



# AmgX Wrapper

A wrapper for quickly coupling AmgX into existing PETSc-based software

# AmgX Wrapper: Make Life Easier

## Declare and initialize a solver

```
AmgXWrapper    solver;  
solver.initialize(communicator & config file);
```



## Bind the matrix A

```
solver.setA(A);
```



## In time-marching loop

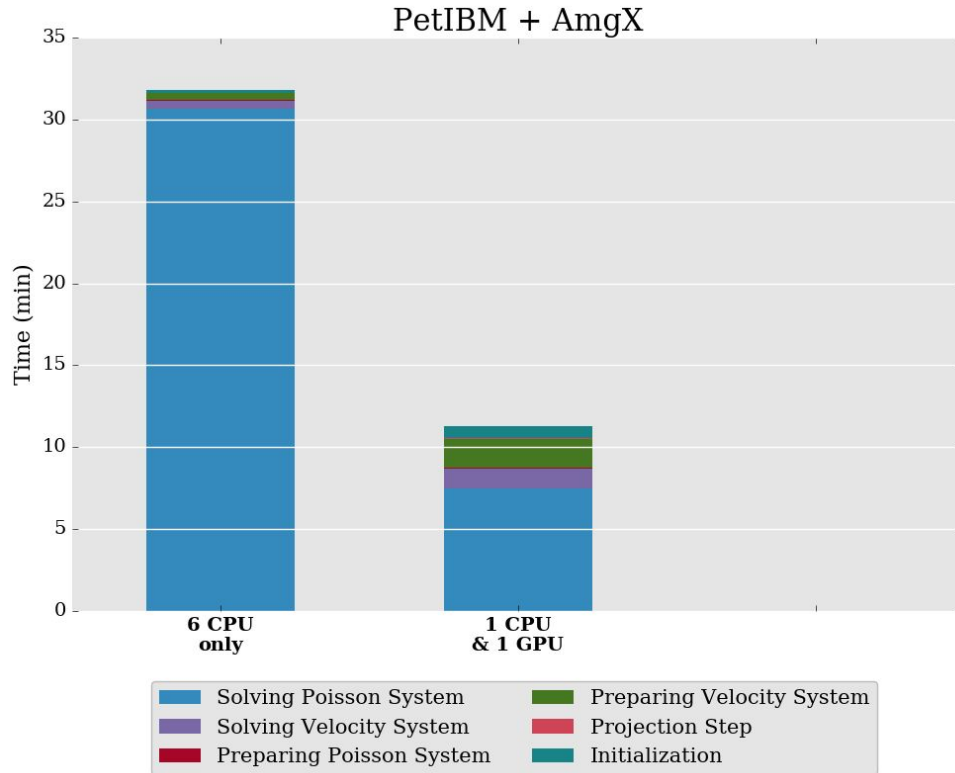
```
solver.solve(x, rhs);
```



## Finalization

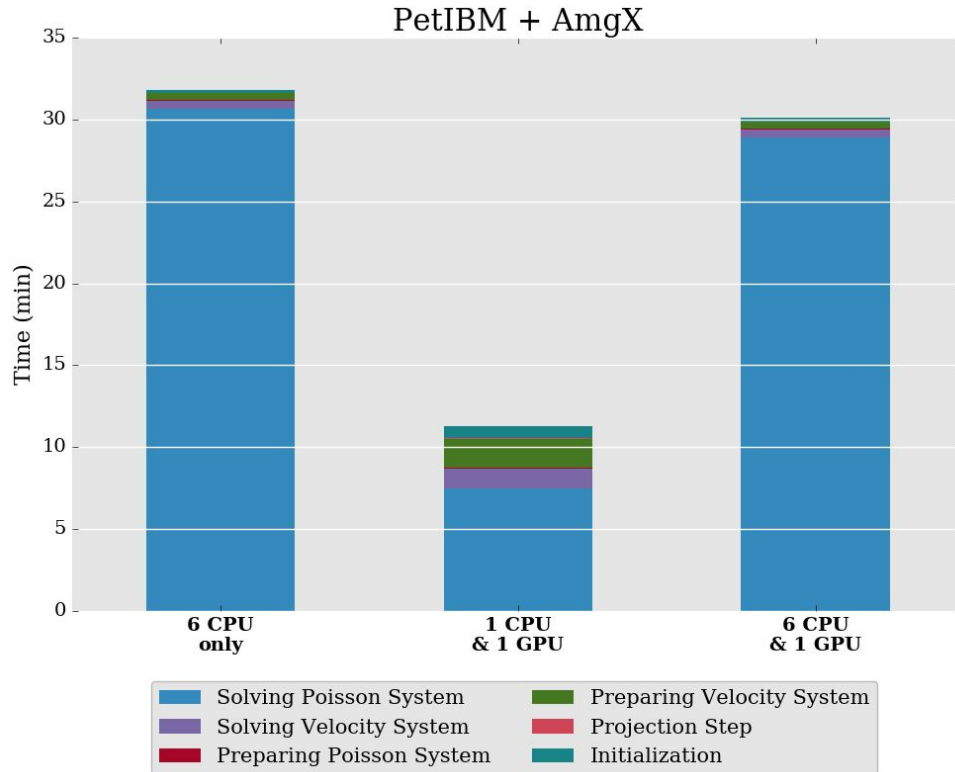
```
solver.finalize();
```

# Example: 2D Cylinder Flow, $Re=40$



- Mesh Size: 2.25M
- 1 NVIDIA K40c
- Velocity:
  - PETSc KSP - CG
  - Block Jacobi
- Modified Poisson
  - AmgX - CG
  - Aggregation AMG

# Example: 2D Cylinder Flow, Re=40



- Mesh Size: 2.25M
- 1 NVIDIA K40c
- Velocity:
  - PETSc KSP - CG
  - Block Jacobi
- Modified Poisson
  - AmgX - CG
  - Aggregation AMG

# Solution

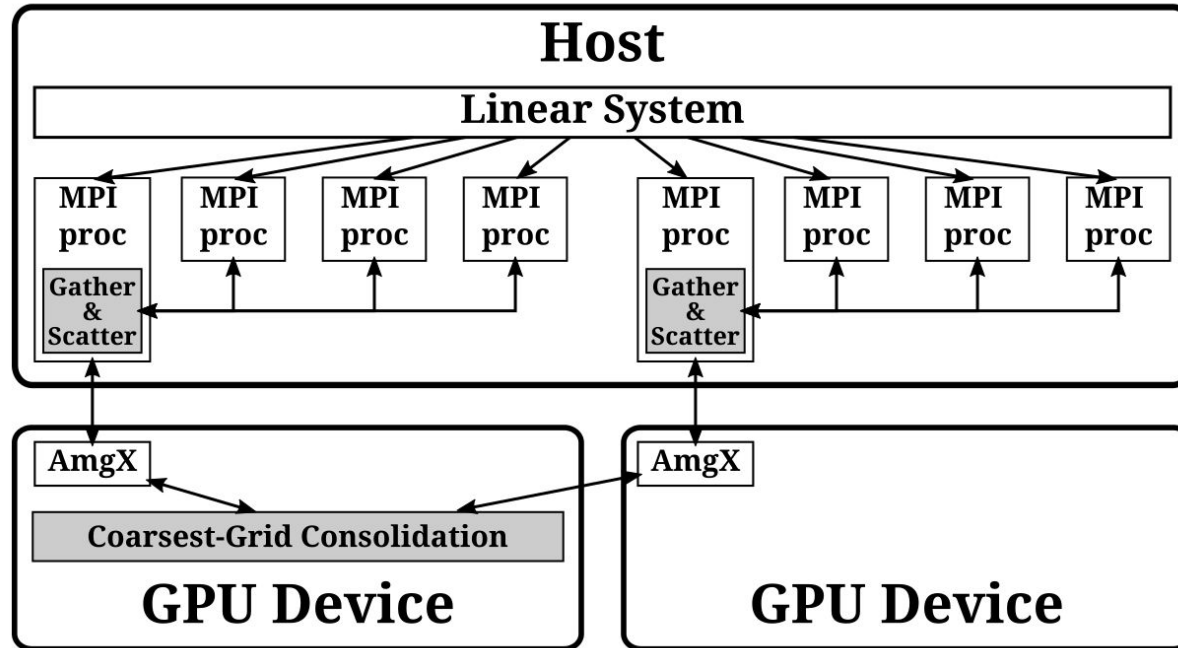
Assure there's always only one subdomain solver on every GPU

# **We want to make using AmgX easy**

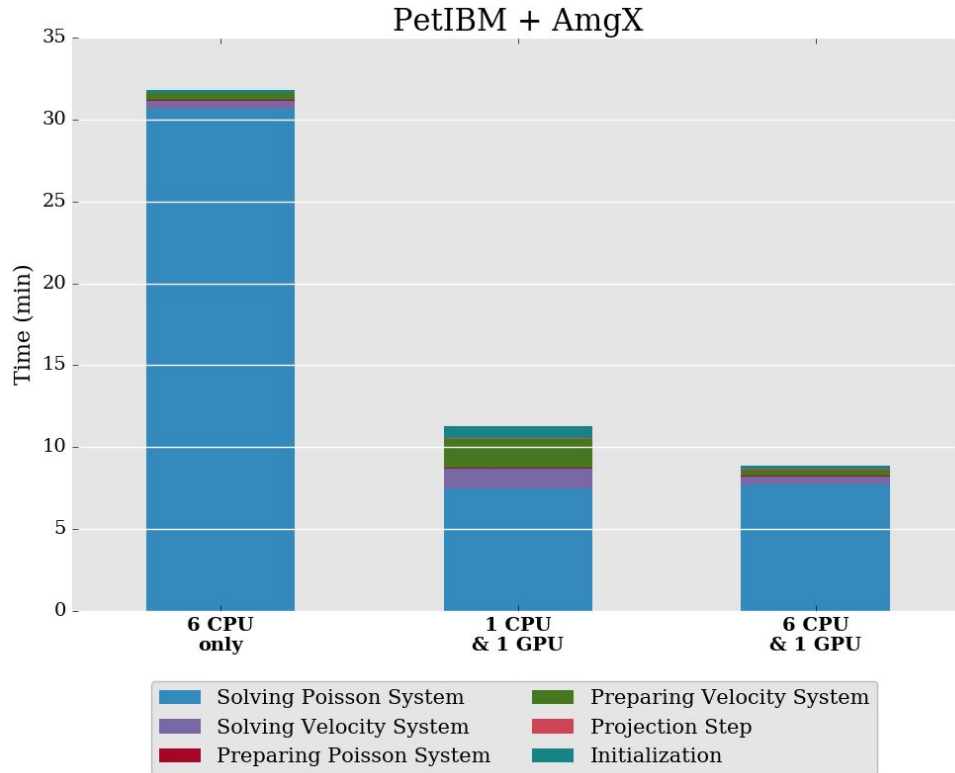
The solution should be implemented in the wrapper, not in PetIBM

# The wrapper makes things easier

No need to modify original codes in PETSc-based applications



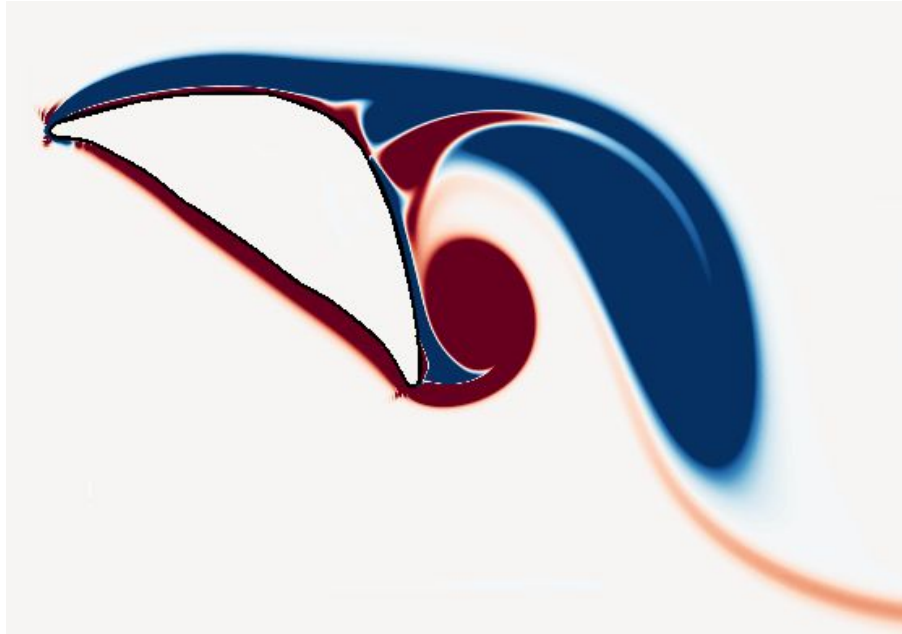
# Back to Example: 2D Cylinder Flow, Re=40



- Mesh Size: 2.25M
- 1 NVIDIA K40c
- Velocity:
  - PETSc KSP - CG
  - Block Jacobi
- Modified Poisson
  - AmgX - CG
  - Aggregation AMG
- AmgX Wrapper



# Benchmark: Flying Snakes

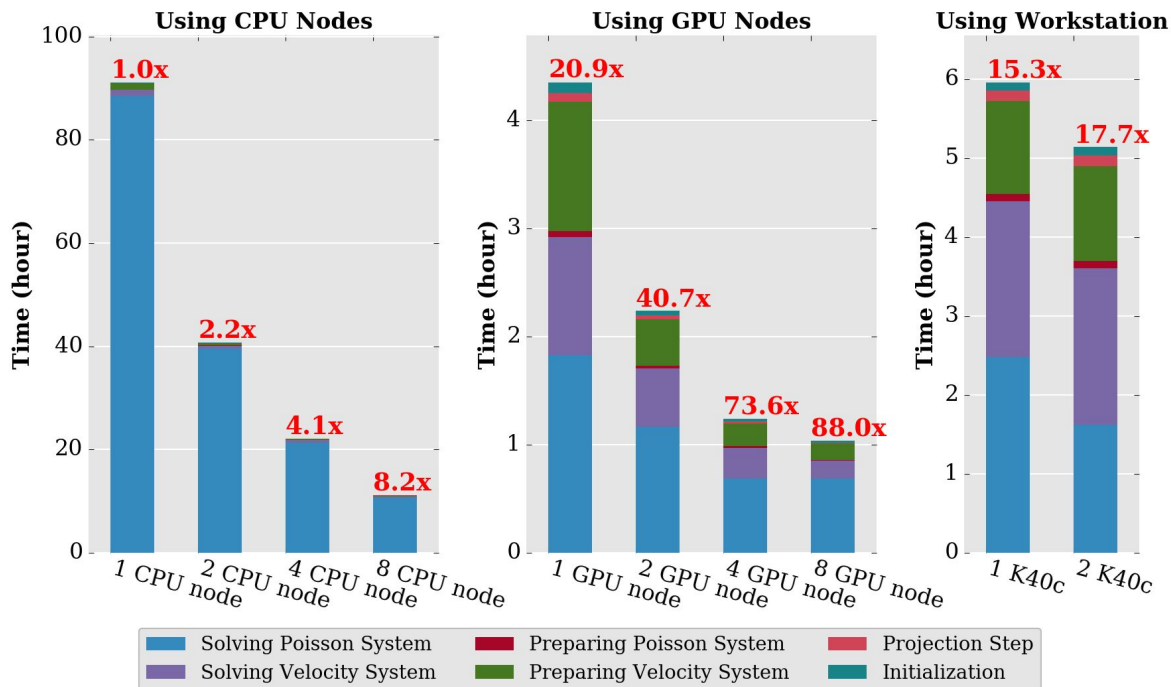


- **Anush Krishnan et. al. (2014)<sup>†</sup>**
  - **Re=2000**
  - **AoA=35**
  - **Mesh Size: 2.9M**

<sup>†</sup>A. Krishnan, J. Socha, P. Vlachos and L. Barba, "Lift and wakes of flying snakes", *Physics of Fluids*, vol. 26, no. 3, p. 031901, 2014.

# Example: Flying Snakes

Flying Snake, 2D, Re=2000, AoA=35



- **Per CPU node:**

- 2 Intel E5-2620 (12 cores)

- **Per GPU node:**

- 1 CPU node (12 cores)
- 2 NVIDIA K20

- **Workstation:**

- Intel i7-5930K (6 cores)
- 1 or 2 K40c

**Time is money**

# Potential Savings and Benefits: Hardware

For our application, enabling multi-GPU computing reduces

- costs on extra hardware,
  - motherboards, memory, hard drives, cooling systems, power supplies, Infiniband switches, physical space... etc.
- works and human resources on managing clusters,
- socket to socket communications
- potential runtime crash due to single node failure or network failure, and
- time spent on queue at any HPC centers

# Potential saving on cloud HPC service

Running GPU-enabled CFD applications with cloud HPC service may save  
a lot

# Potential Saving and Benefits: Cloud HPC Service

Reduce execution time and needed nodes. For example, on Amazon EC2:

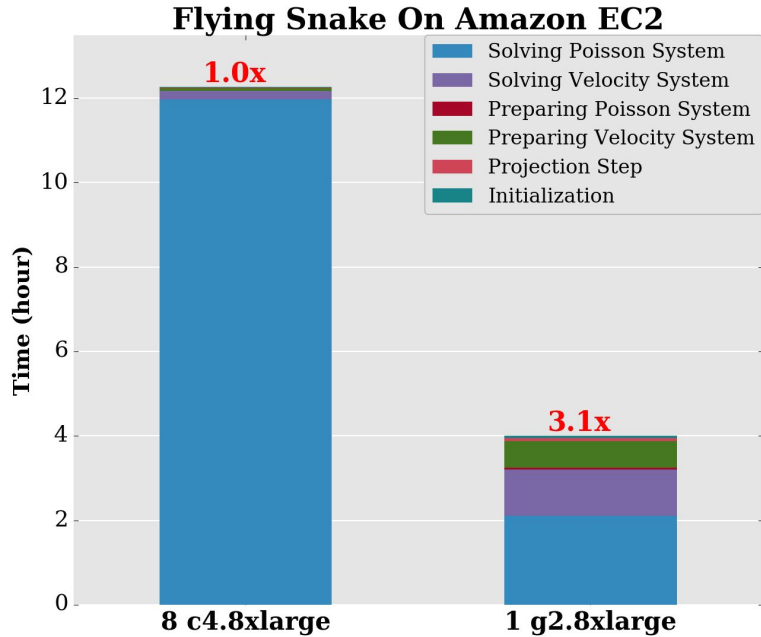
- **GPU nodes - g2.8xlarge:**

- 32 vCPU (Intel E5-2670) + 4 GPUs (Kepler GK104 )
- Official Price: \$2.6 / hr
- Possible Lower Price (Spot Instances): < \$0.75 / hr

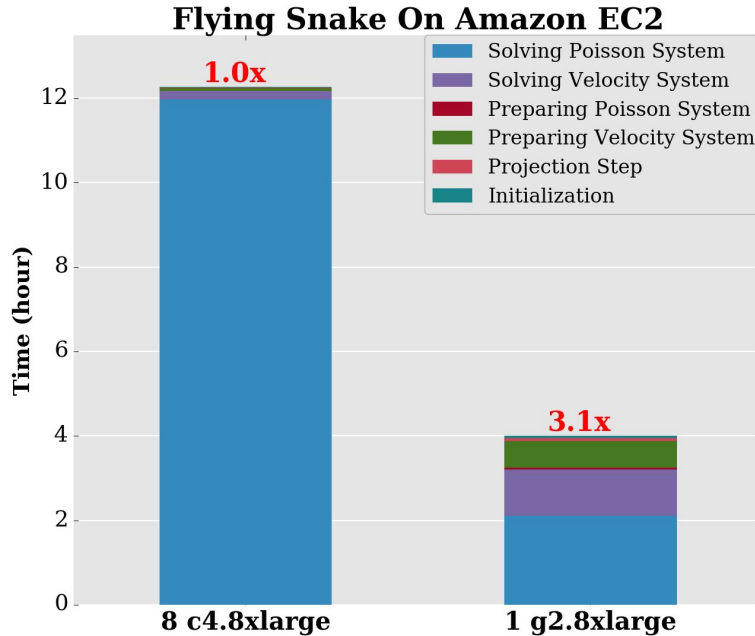
- **CPU nodes - c4.8xlarge**

- 36 vCPU (Intel E5-2663)
- Official Price: \$1.675 / hr
- Possible Lower Price (Spot Instances): < \$0.6 / hr

# Potential Saving and Benefits: Cloud HPC Service



# Potential Saving and Benefits: Cloud HPC Service



- **CPU:**

$$12.5 \text{ hr} \times \$1.675 / \text{hr} \times 8 \text{ nodes} = \text{\textcolor{red}{\$167.5}}$$

- **GPU:**

$$4 \text{ hr} \times \$2.6 / \text{hr} \times 1 \text{ node} = \text{\textcolor{red}{\$10.4}}$$



# Conclusion

- **AmgX and our wrapper**
  - <https://developer.nvidia.com/amgx>
  - <https://github.com/barbagroup/AmgXWrapper>
- **PetIBM with AmgX enabled:**
  - <https://github.com/barbagroup/PetIBM/tree/AmgXSolvers>
- **Speed up in a real application: flying snake**
- **Time is money**
- **Complete technical paper:**
  - <http://goo.gl/0DM1Vw>

# Thanks!

## Acknowledgement:

Dr. Joe Eaton, NVIDIA

## Technical paper:

<http://goo.gl/0DM1Vw>

## Contact us:

Website:

<http://lorenabarba.com/>

GitHub:

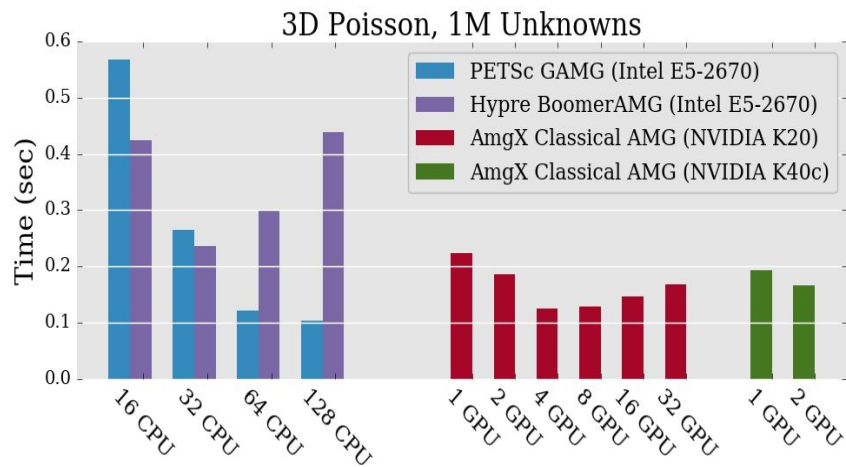
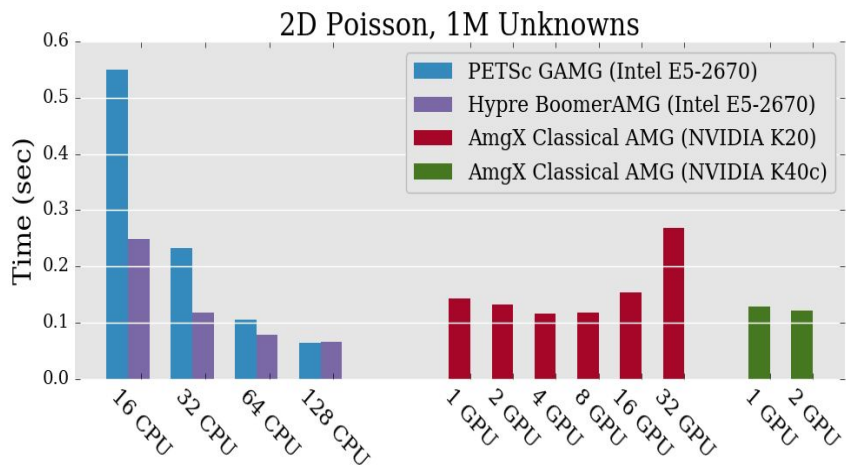
<https://github.com/barbagroup/>



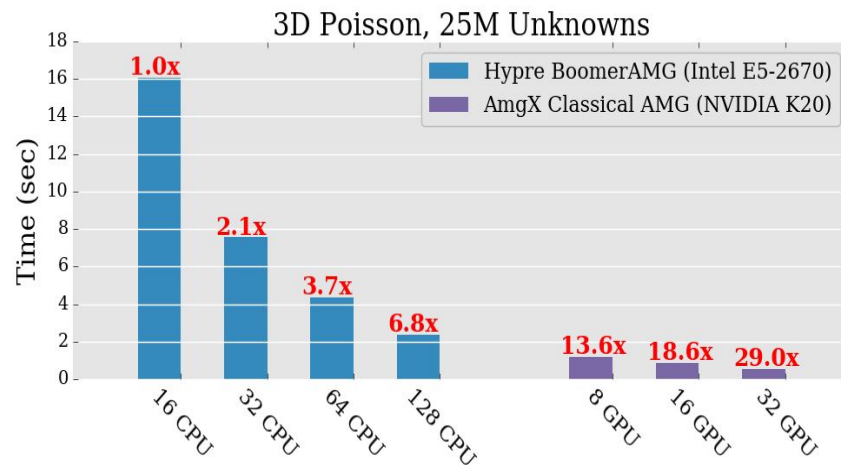
# Q & A

# Extra Slides

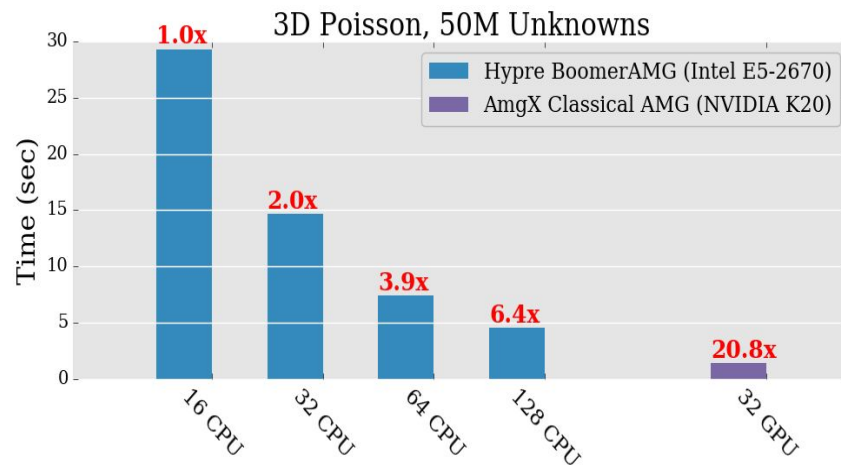
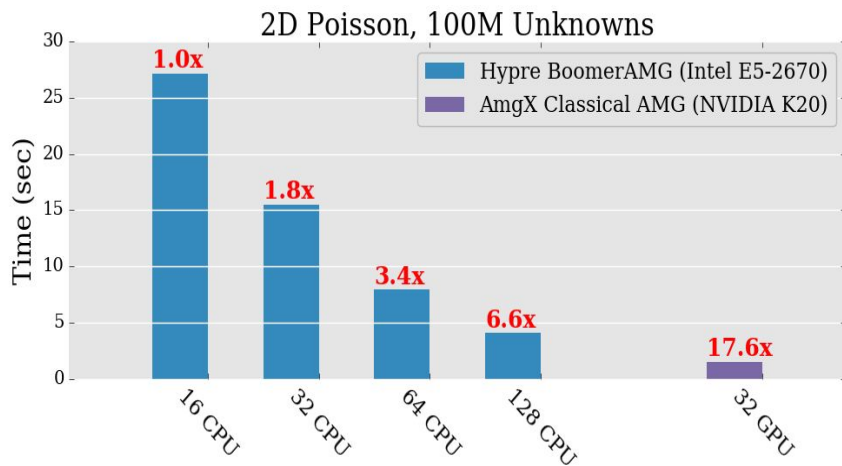
# Example: Small-Size Problems



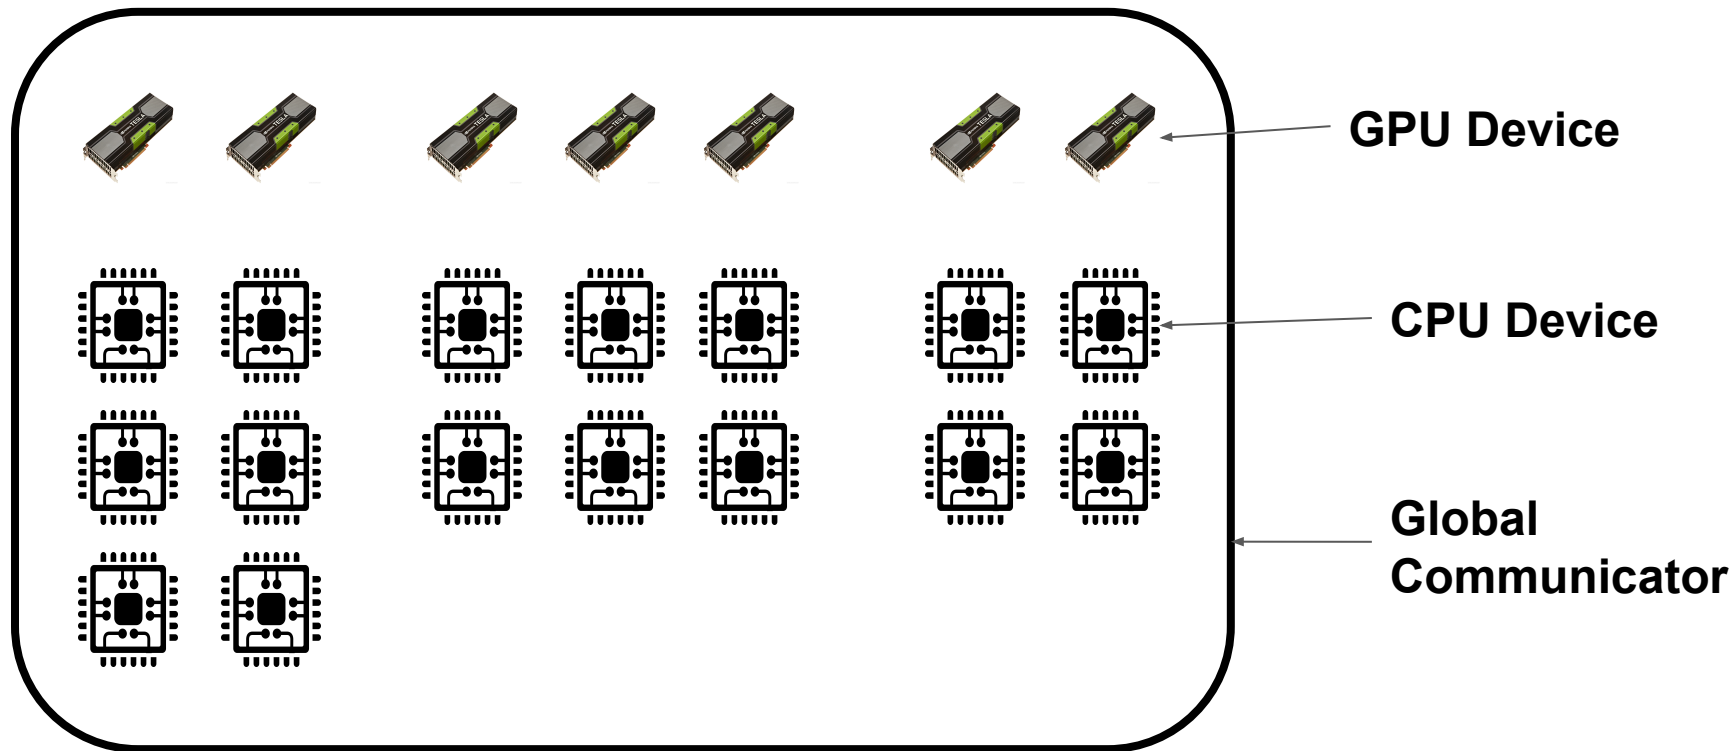
# Example: Medium-Size Problems



# Example: Large-Size Problems

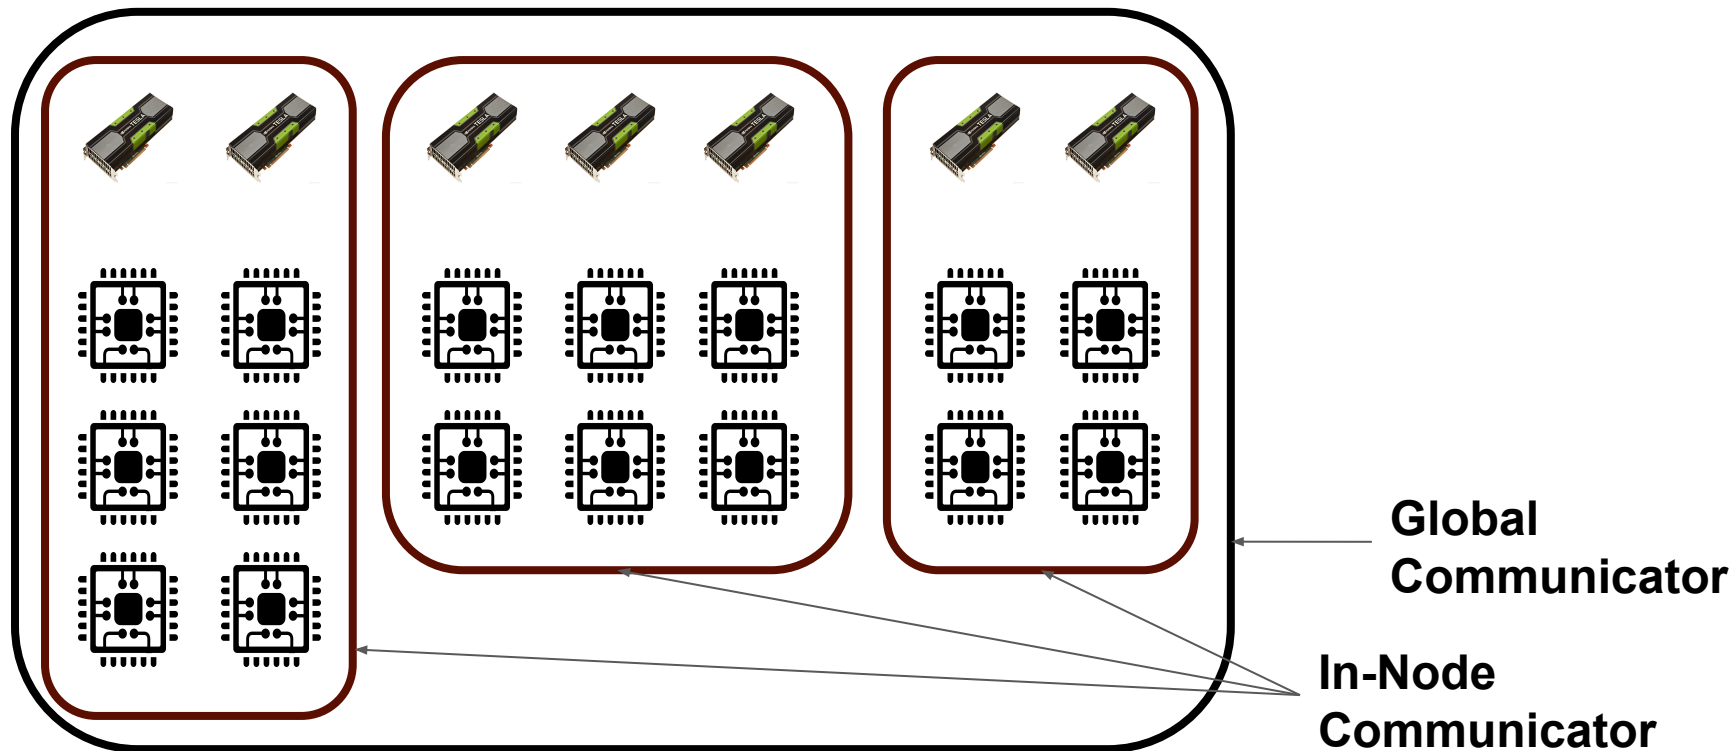


# Our AmgX Wrapper handle this case !

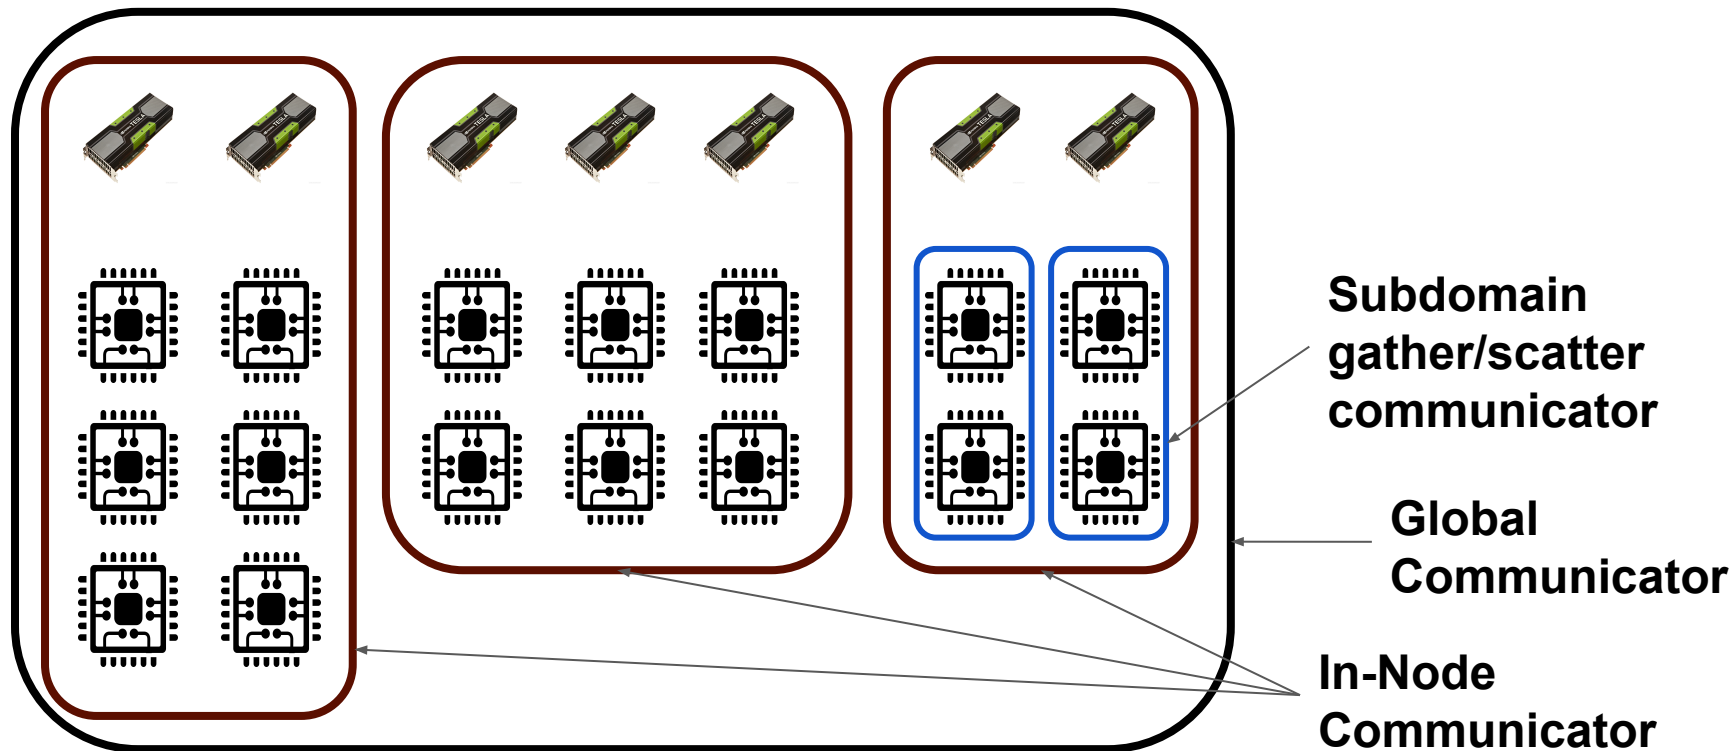




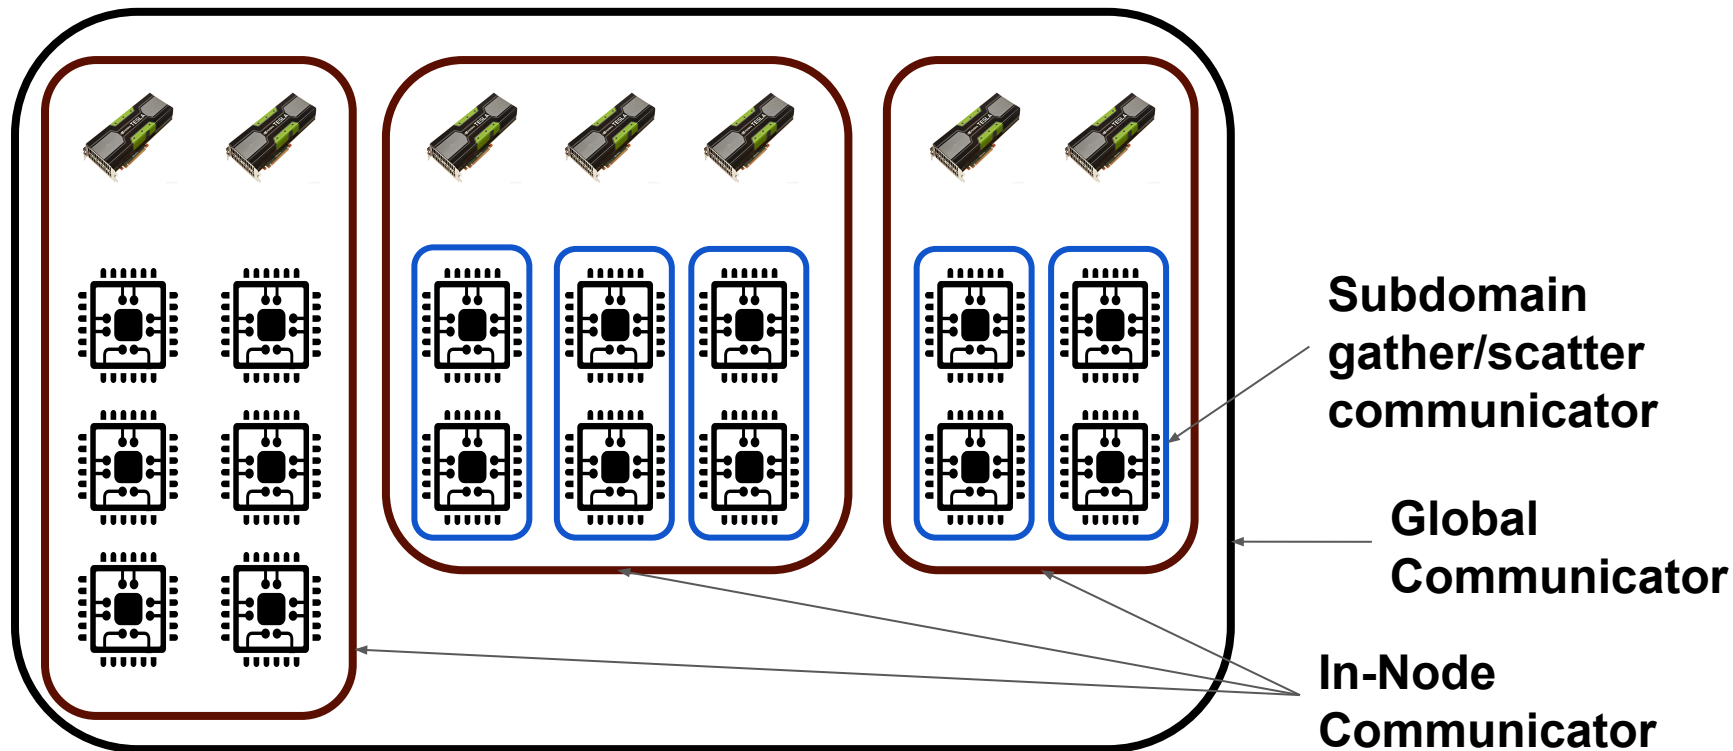
# Our AmgX Wrapper handle this case !



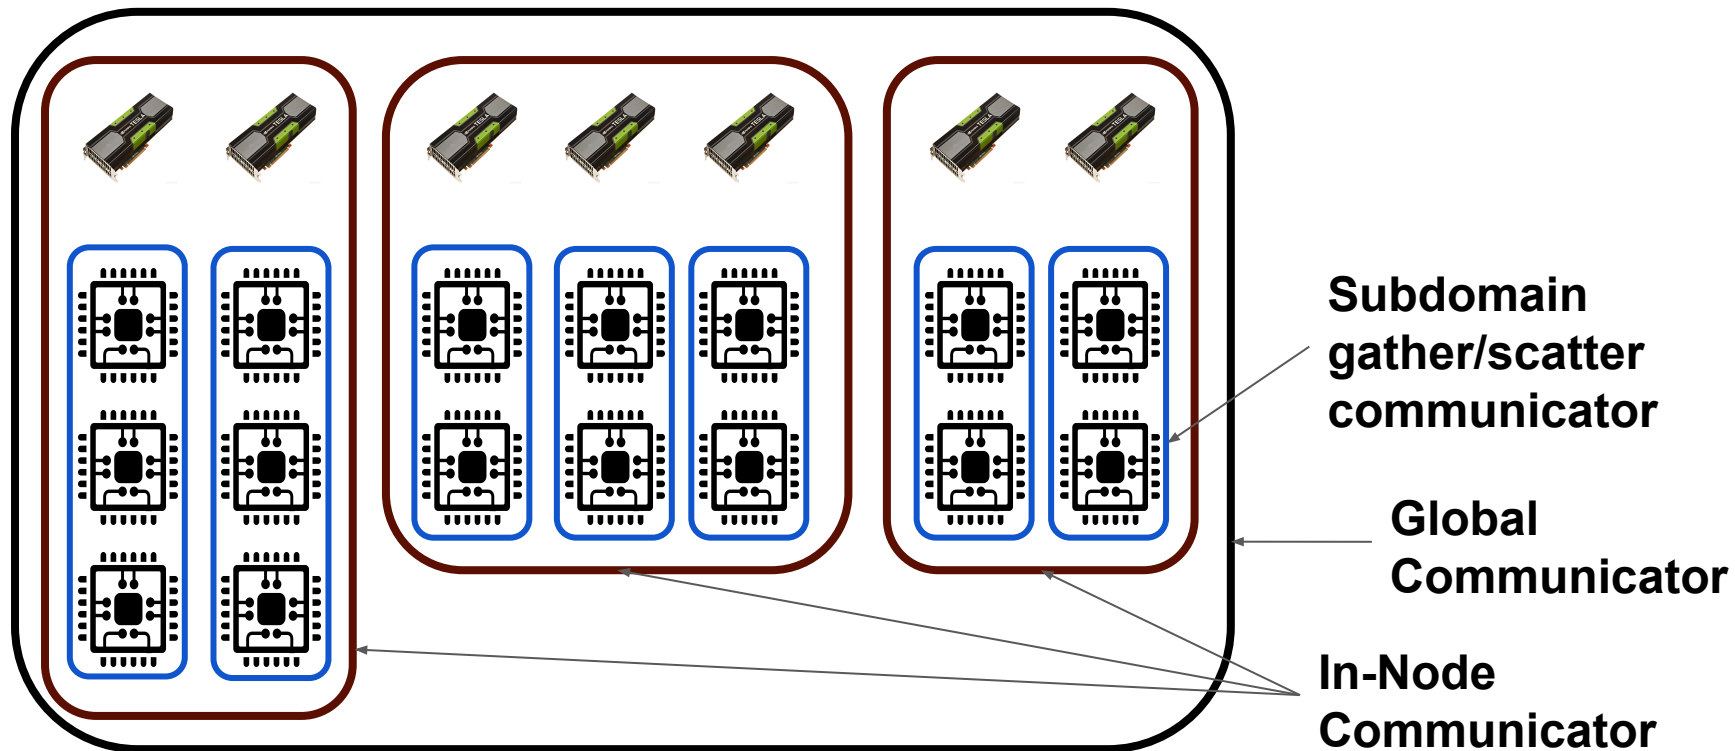
# Our AmgX Wrapper handle this case !



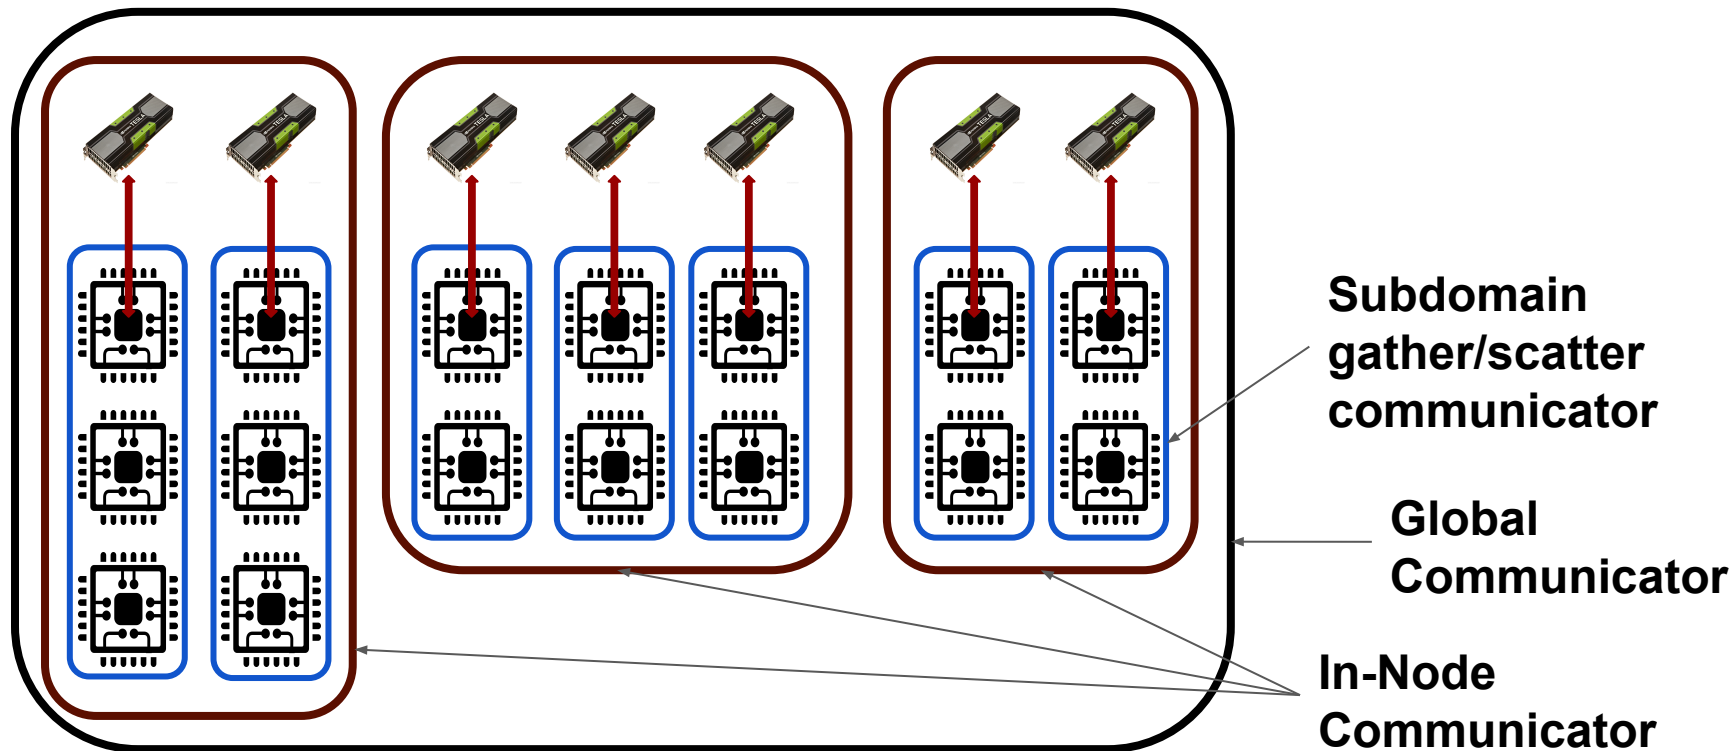
# Our AmgX Wrapper handle this case !



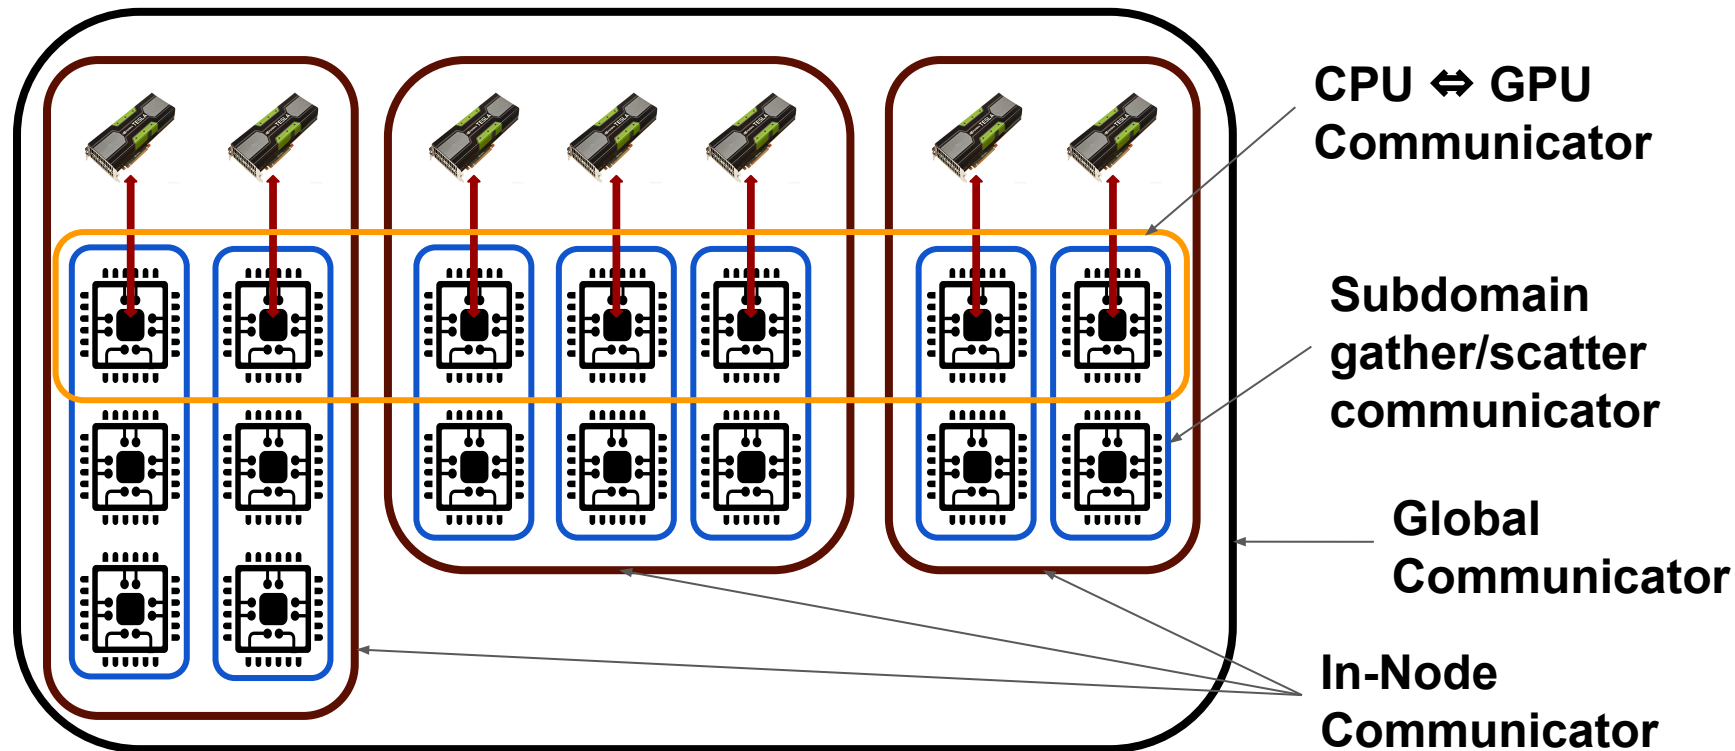
# Our AmgX Wrapper handle this case !



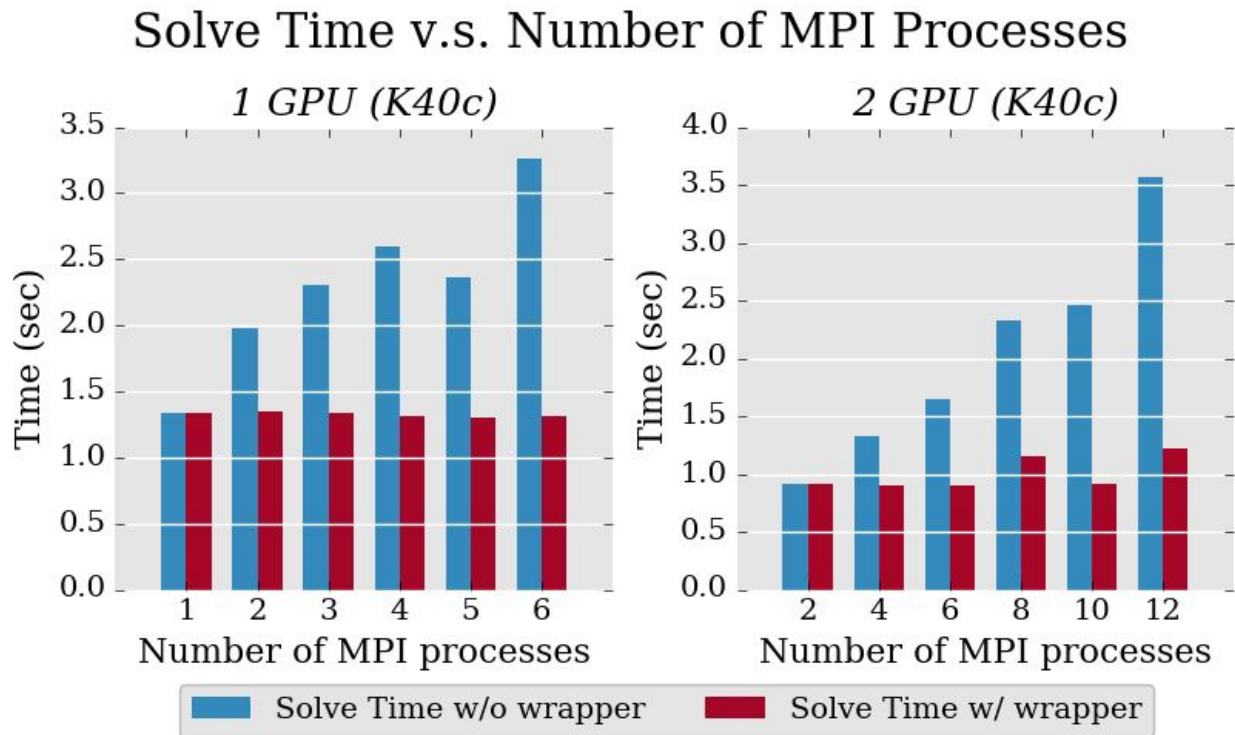
# Our AmgX Wrapper handle this case !



# Our AmgX Wrapper handle this case !

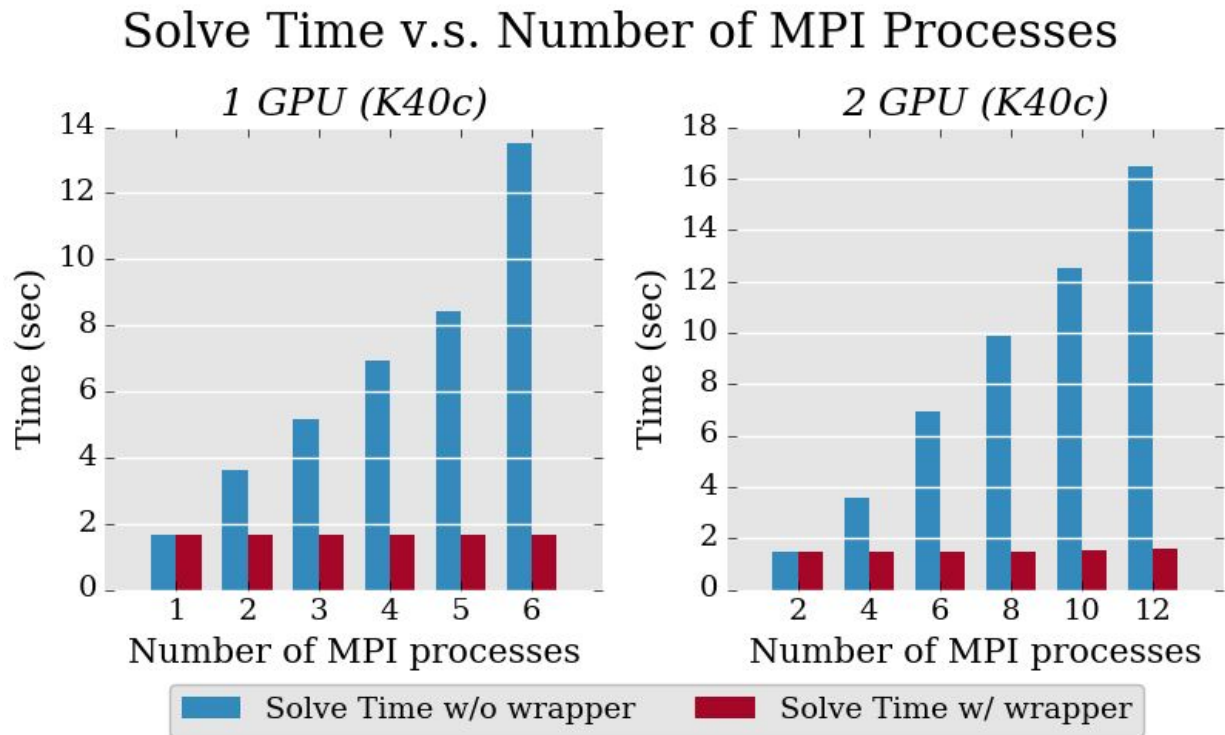


# Check: 3D Poisson



- 6M unknowns
- Solver:
  - CG
  - Classical AMG

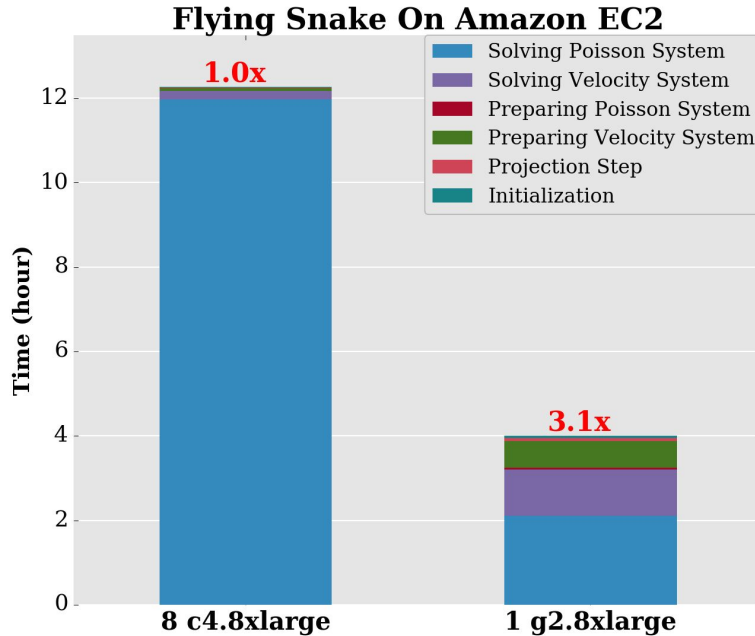
# Check: Modified Poisson Equation



- 2D Cylinder, Re 40
- 2.25M unknowns
- Solver:
  - CG
  - Aggregation AMG



# Potential Saving and Benefits: Cloud HPC Service



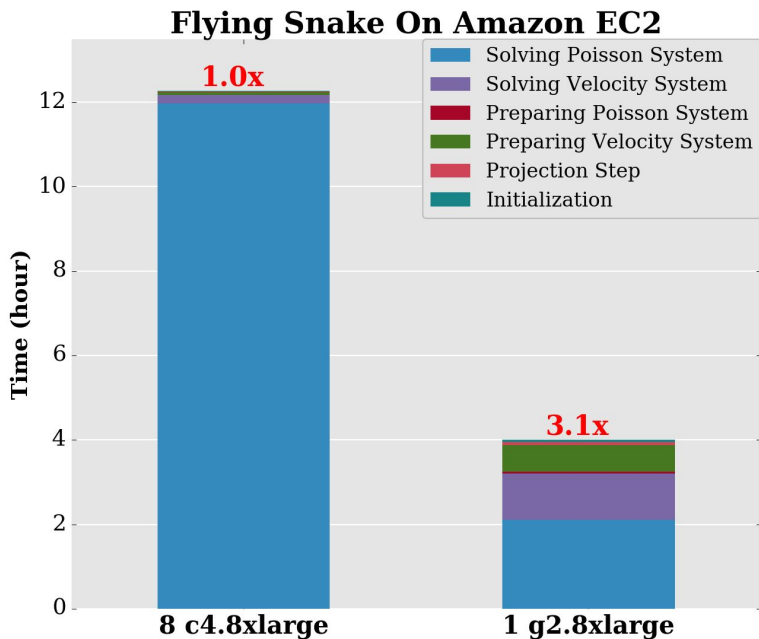
- **Using Spot Instances**

- **CPU:**

- $12.5 \text{ hr} \times \$0.5^{\dagger} / \text{hr} \times 8 \text{ nodes} = \text{\textcolor{red}{\$50.0}}$

<sup>†</sup>This is the prices of the spot instances we used at that time.

# Potential Saving and Benefits: Cloud HPC Service



- **Using Spot Instances**

- **CPU:**

- $12.5 \text{ hr} \times \$0.5^{\dagger} / \text{hr} \times 8 \text{ nodes} = \underline{\$50.0}$

- **GPU:**

- $4 \text{ hr} \times \$0.5^{\dagger} / \text{hr} \times 1 \text{ node} = \underline{\$2.0}$

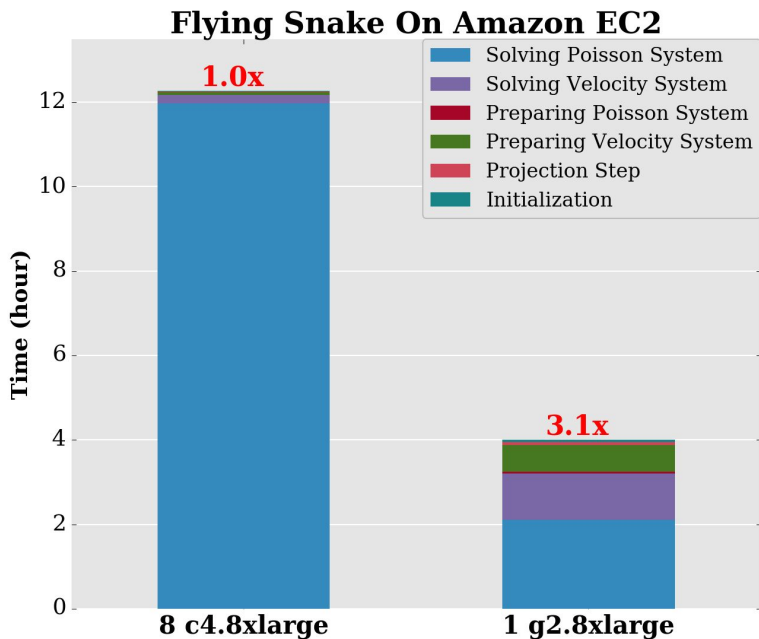
- **Using Official Price:**

- **CPU:**

- $12.5 \text{ hr} \times \$1.675 / \text{hr} \times 8 \text{ nodes} = \underline{\$167.5}$

<sup>†</sup>This is the prices of the spot instances we used at that time.

# Potential Saving and Benefits: Cloud HPC Service



- **Using Spot Instances**

- **CPU:**

- $12.5 \text{ hr} \times \$0.5^{\dagger} / \text{hr} \times 8 \text{ nodes} = \underline{\$50.0}$

- **GPU:**

- $4 \text{ hr} \times \$0.5^{\dagger} / \text{hr} \times 1 \text{ node} = \underline{\$2.0}$

- **Using Official Price:**

- **CPU:**

- $12.5 \text{ hr} \times \$1.675 / \text{hr} \times 8 \text{ nodes} = \underline{\$167.5}$

- **GPU:**

- $4 \text{ hr} \times \$2.6 / \text{hr} \times 1 \text{ node} = \underline{\$10.4}$

<sup>†</sup>This is the prices of the spot instances we used at that time.

# PetIBM

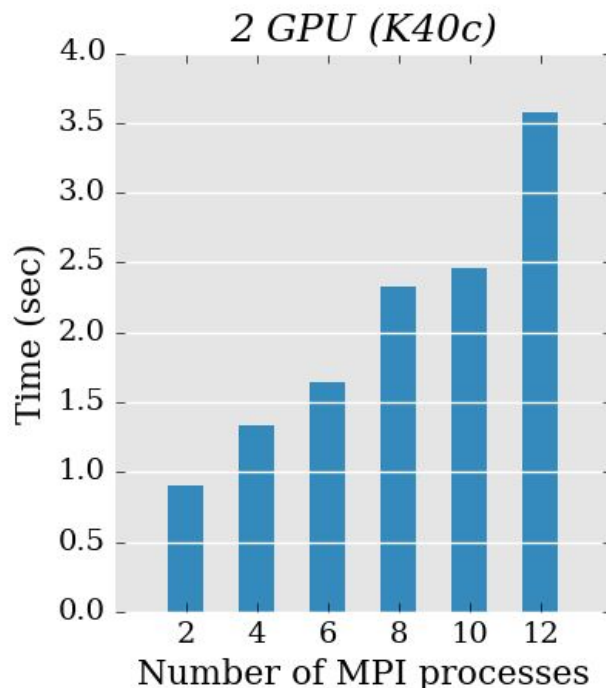
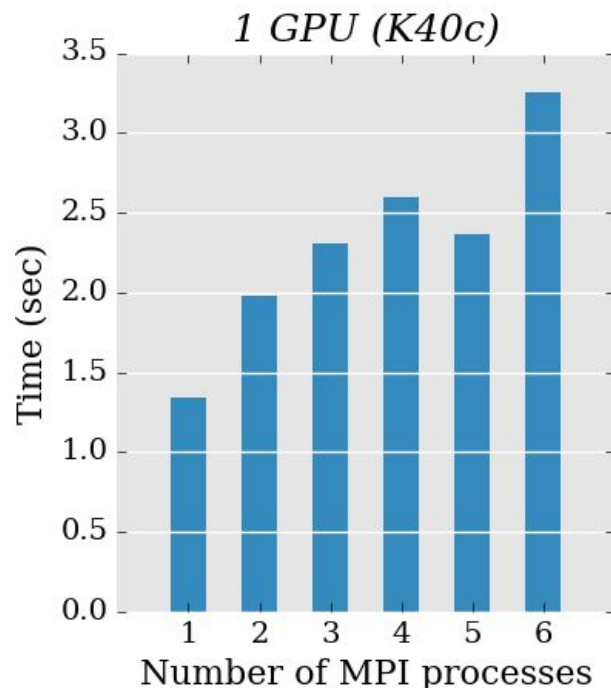
Solving Poisson systems in CFD solvers is already tough, but ...

# AmgX

- 
- **C API**
- **Unified Virtual Addressing**
- **Smoothers:**
  - Block-Jacobi, Gauss-Seidel, incomplete LU, Polynomial, dense LU ... etc
- **Cycles:**
  - V, W, F, CG, CGF

# Tests: 3D Poisson

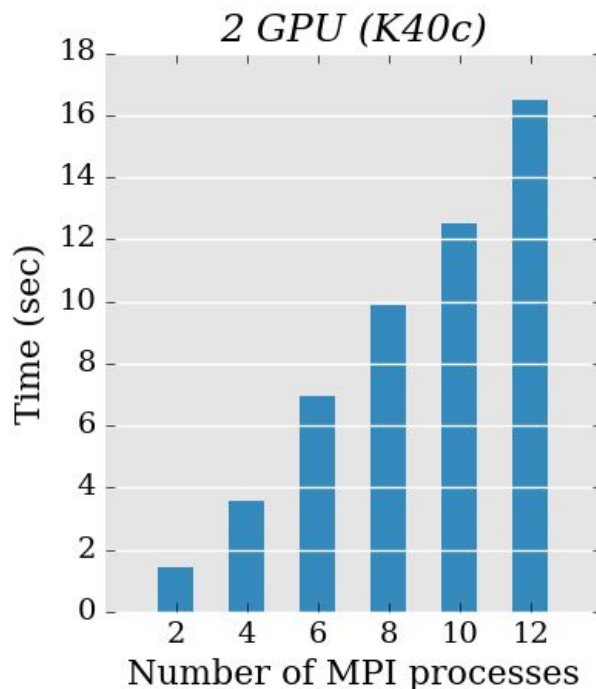
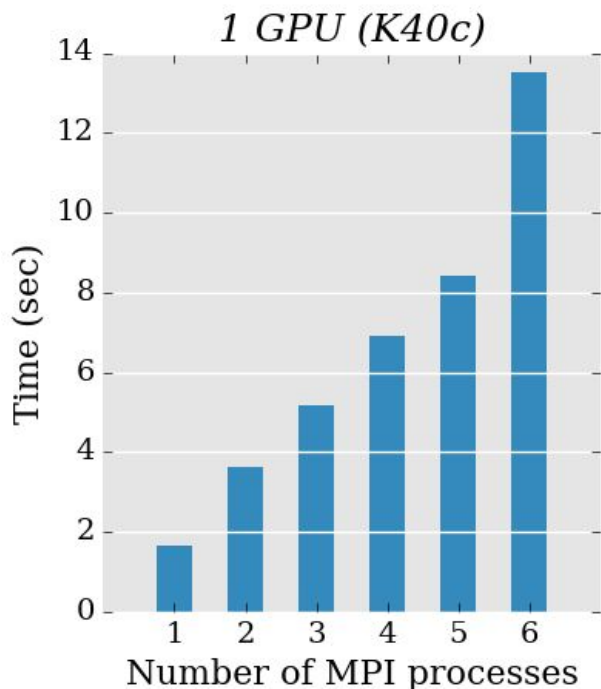
Solve Time v.s. Number of MPI Processes



- 6M unknowns
- Solver:
  - CG
  - Classical AMG

# Tests: Modified Poisson Equation

Solve Time v.s. Number of MPI Processes



- 2D Cylinder, Re 40
- 2.25M unknowns
- Solver:
  - CG
  - Aggregation AMG